

Session 4 - ggplot2

Alex Mounsey

07/10/2020

Exercise 1

Consider again the data collected by a farm management company from nine farms in Devon about the number of cattle and sheep. It turns out that other information is available, including the farm location and whether the opportunity to participate in a working farm holiday is offered:

$$\begin{array}{l} \text{Cattle} = \begin{pmatrix} 348 \\ 407 \\ 1064 \\ 750 \\ 593 \\ 1867 \\ 471 \\ 935 \\ 1443 \end{pmatrix} \quad \text{Sheep} = \begin{pmatrix} 110 \\ 179 \\ 303 \\ 173 \\ 182 \\ 458 \\ 151 \\ 140 \\ 222 \end{pmatrix} \quad \text{Location} = \begin{pmatrix} \text{North} \\ \text{South} \\ \text{South} \\ \text{North} \\ \text{South} \\ \text{North} \\ \text{North} \\ \text{North} \\ \text{South} \end{pmatrix} \quad \text{Holidays} = \begin{pmatrix} \text{No} \\ \text{No} \\ \text{Yes} \\ \text{Yes} \\ \text{No} \\ \text{Yes} \\ \text{No} \\ \text{No} \\ \text{Yes} \end{pmatrix} \end{array}$$

```
cattle <- c(348, 407, 1064, 750, 593, 1867, 471, 935, 1443)
sheep <- c(110, 179, 303, 173, 182, 458, 151, 140, 222)
location <- c('North', 'South', 'South', 'North', 'South', 'North', 'North',
              'North', 'South')
holidays <- c('No', 'No', 'Yes', 'Yes', 'No', 'Yes', 'No', 'No', 'Yes')
```

Define categorical variables `location` and `holidays` as factors and create a dataframe including these newly defined variables, as well as the `cattle` and `sheep` variables that you've worked with before:

```
location_f <- factor(location, levels = c('North', 'East', 'South', 'West'))
holidays_f <- factor(holidays, levels = c('Yes', 'No'))

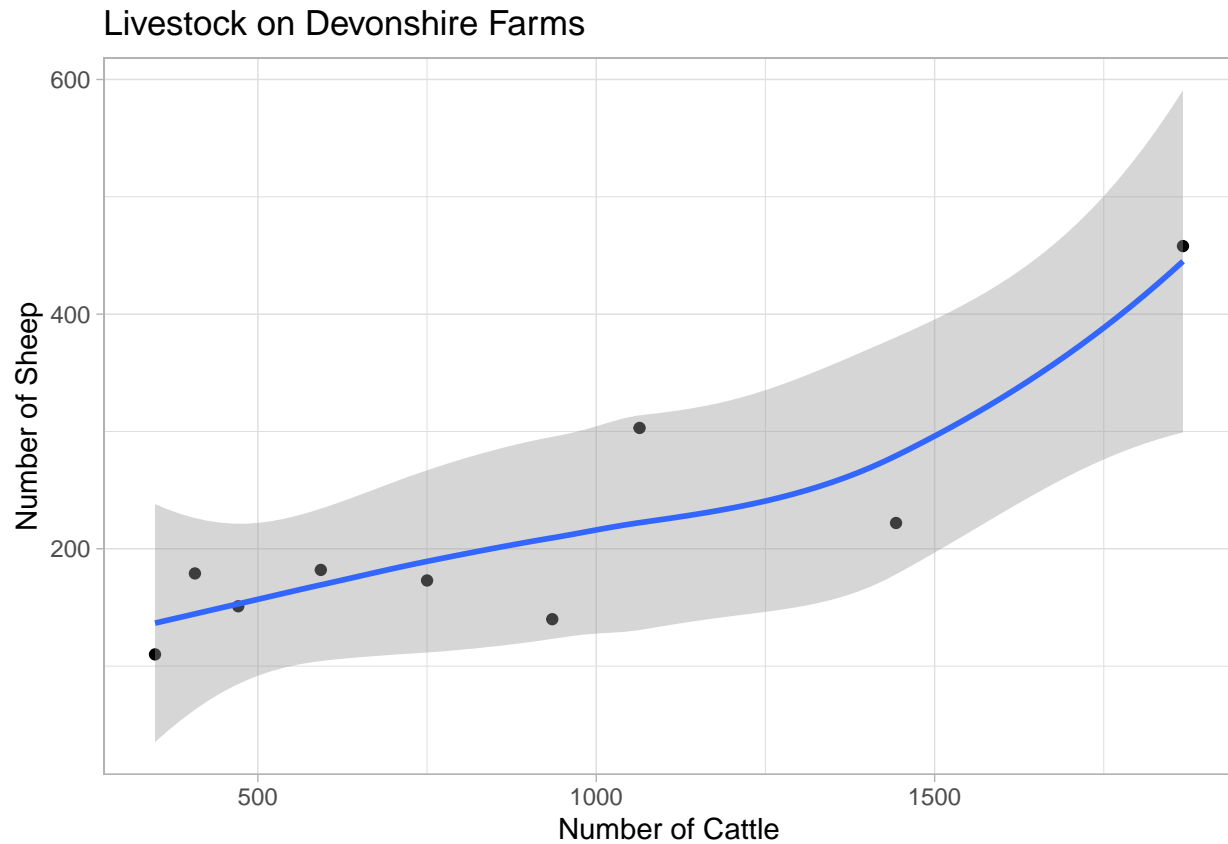
farms_df <- data.frame(cattle, sheep, location_f, holidays_f) %>%
  rename(location = location_f, holidays = holidays_f)

head(farms_df)
```

```
##   cattle sheep location holidays
## 1    348   110   North        No
## 2    407   179   South        No
## 3   1064   303   South        Yes
## 4    750   173   North        Yes
## 5    593   182   South        No
## 6   1867   458   North        Yes
```

Reproduce the plots shown in the tutorial file:

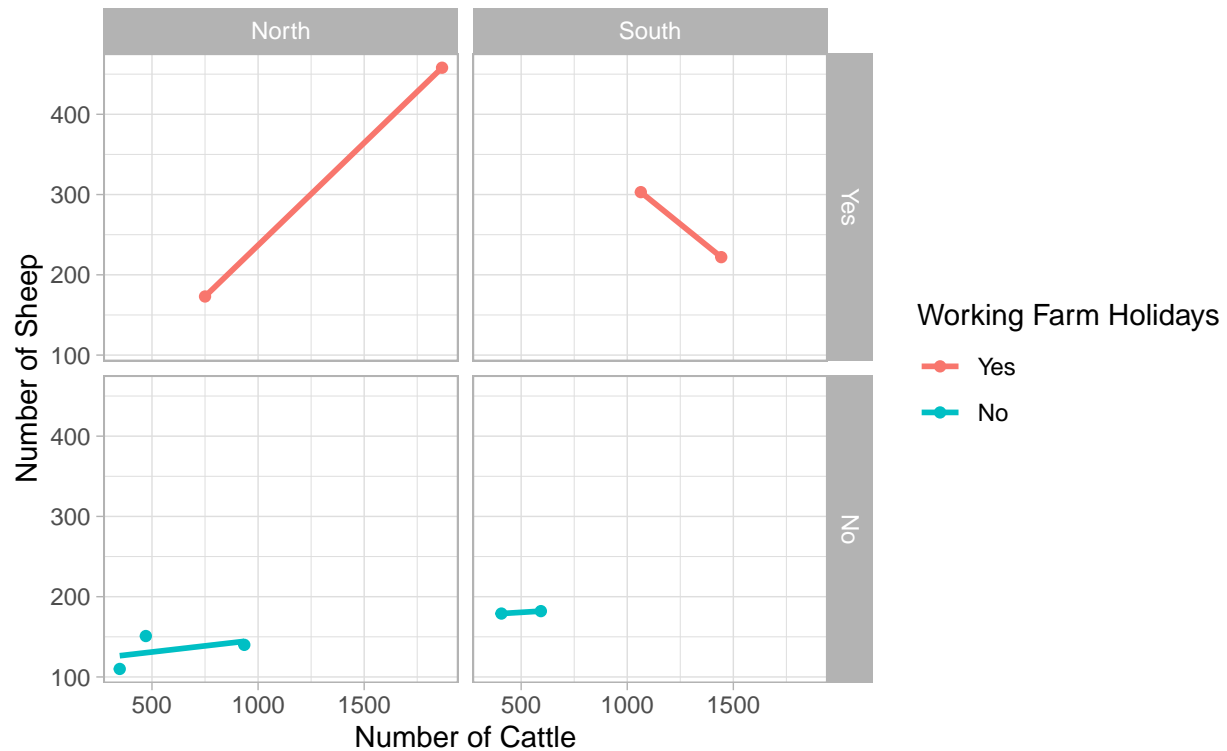
```
ggplot(farms_df, aes(x = cattle, y = sheep)) +  
  theme_light() + geom_point() + geom_smooth(span = 1) +  
  labs(x = "Number of Cattle", y = "Number of Sheep",  
       title = "Livestock on Devonshire Farms")
```



```
ggplot(farms_df, aes(x = cattle, y = sheep, colour = holidays)) +  
  theme_light() + geom_point() + geom_smooth(method = 'lm', se = F) +  
  facet_grid(holidays ~ location) +  
  labs(x = "Number of Cattle", y = "Number of Sheep",  
       colour = "Working Farm Holidays",  
       title = "Livestock on Devonshire Farms",  
       subtitle = "Faceted by location and working farm holidays")
```

Livestock on Devonshire Farms

Faceted by location and working farm holidays



Exercise 2

Read the questionnaire data into a dataframe, `df`, in the usual way:

```
q_data <- read_csv('../data/MATH513_Questionnaire_Data.csv')
head(q_data)
```

```
## # A tibble: 6 x 19
##   Height  Age Sex  BirthPlace SiblingsNo EatMeat DrinkCoffee LikeBeer Sports
##   <dbl> <dbl> <chr> <chr>          <dbl> <chr>    <chr>      <chr>    <chr>
## 1   170   23 Fema~ essex           1 Yes      Yes        No      Yes
## 2   188  22.4 Male London           1 Yes      Yes        No      No
## 3   180  30.1 Male Athens          0 Yes      Yes        Yes     Yes
## 4   185   21 Male China            0 Yes      Yes        Yes     Yes
## 5   170  22.1 Fema~ Plymouth        2 Yes      Yes        No      No
## 6   182   25 Male Nigeria          4 Yes      No         No      Yes
## # ... with 10 more variables: Driver <chr>, LeftHanded <chr>, Abroad <chr>,
## #   Sleep <dbl>, Rent <dbl>, Happy_accommodation <chr>, Distance <dbl>,
## #   Travel_time <dbl>, Mode_of_transport <chr>, Safe <chr>
```

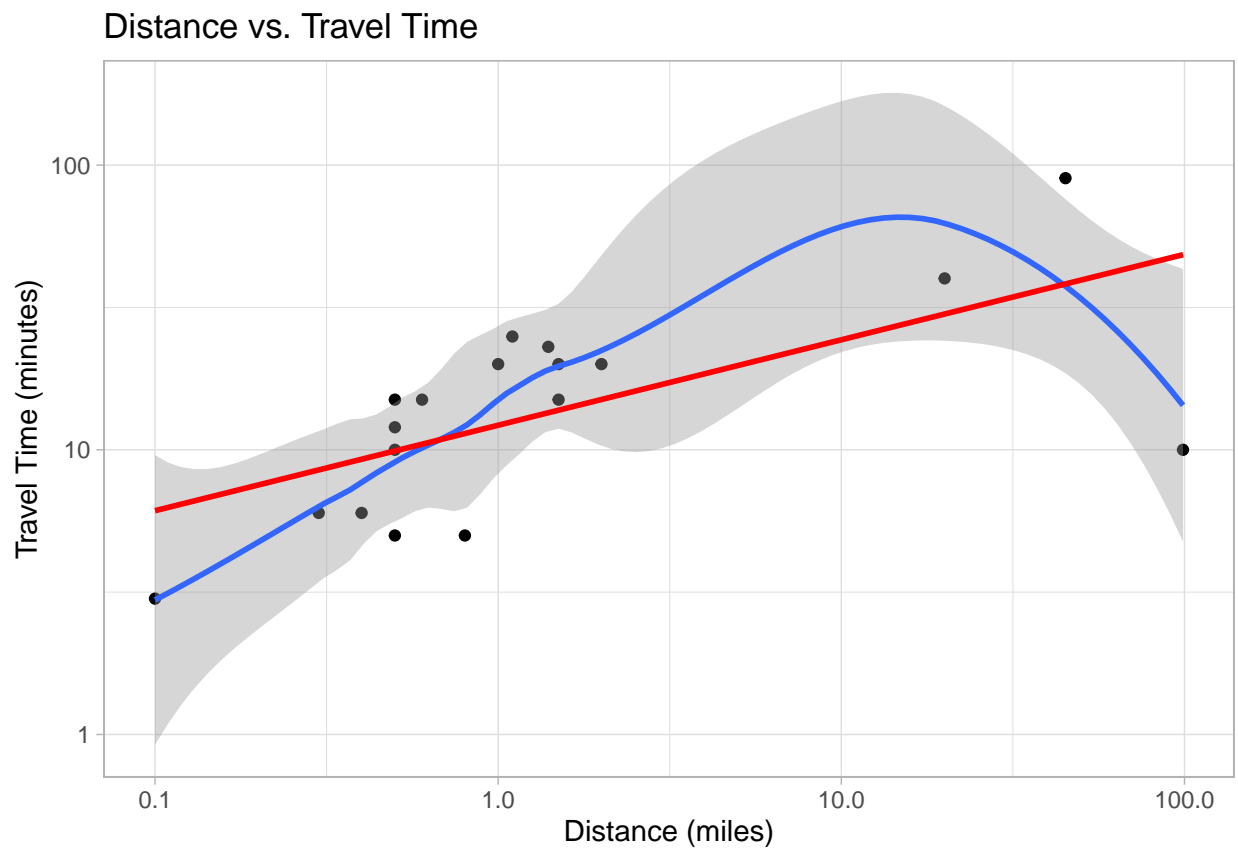
Using the `dplyr` package, select and display the columns (*variables*) `Travel_time` and `Distance`:

```
q_data %>%
  select(Travel_time, Distance) %>%
  head()
```

```
## # A tibble: 6 x 2
##   Travel_time Distance
##       <dbl>     <dbl>
## 1         20         1
## 2          5         0.8
## 3          3         0.1
## 4         20         2
## 5         40        20
## 6         15         0.6
```

Your goal is to understand how Travel_time is dependent on Distance. Use ggplot2 to plot Travel_time against Distance:

```
ggplot(q_data, aes(x = Distance, y = Travel_time)) +
  theme_light() +
  geom_point() + geom_smooth() + geom_smooth(method = 'lm', se = F, colour = 'red') +
  scale_x_log10() + scale_y_log10() +
  labs(x = "Distance (miles)", y = "Travel Time (minutes)",
       title = "Distance vs. Travel Time")
```



Logarithmic Axis Scales

The logarithmic scale is used in situations where data points range between extremely large values and extremely small values, such that the detail of the smaller values is not lost.

Regression Line

The general mathematical equation of a straight line is:

$$y = \alpha + \beta x$$

where: α is known as the intercept and β is known as the slope.

When we're handling data, we can't expect the data points to lie perfectly on the line, so we allow an error. The simple linear regression model therefore takes the form:

$$y = \alpha + \beta x + error$$

Exercise 3

The file `companies.xlsx` contains information collected from 100 companies belonging to different sectors of the economy. The following variables have been recorded:

- `company`: The company's ID
- `net_income_2015`: The company's net income in 2015
- `net_income_2014`: The company's net income in 2014
- `oper_result`: The company's operational result in 2015
- `lab_cost`: The company's labor expenditure in 2015
- `n_empl`: The number of employees in the company
- `sector`: The economic sector the company operates within

```
c_data <- read_excel('../data/companies.xlsx')
head(c_data)
```

```
## # A tibble: 6 x 7
##   company net_income_2015 net_income_2014 oper_result lab_cost n_empl sector
##   <dbl>         <dbl>         <dbl>         <dbl>         <dbl> <dbl> <dbl>
## 1         1         23109600         24324700        -277747      4859414 103000      4
## 2         2         20083509         17384211         2861590      8370166 109860      4
## 3         3         19453500         16665000         2747100      5288500  89475      4
## 4         4         14308944         14144552          256078       286478   4022      2
## 5         5         11243950          9556867         2025679       391221   6415      3
## 6         6          9350460          9061480        -4356350     10739080 179311      4
```

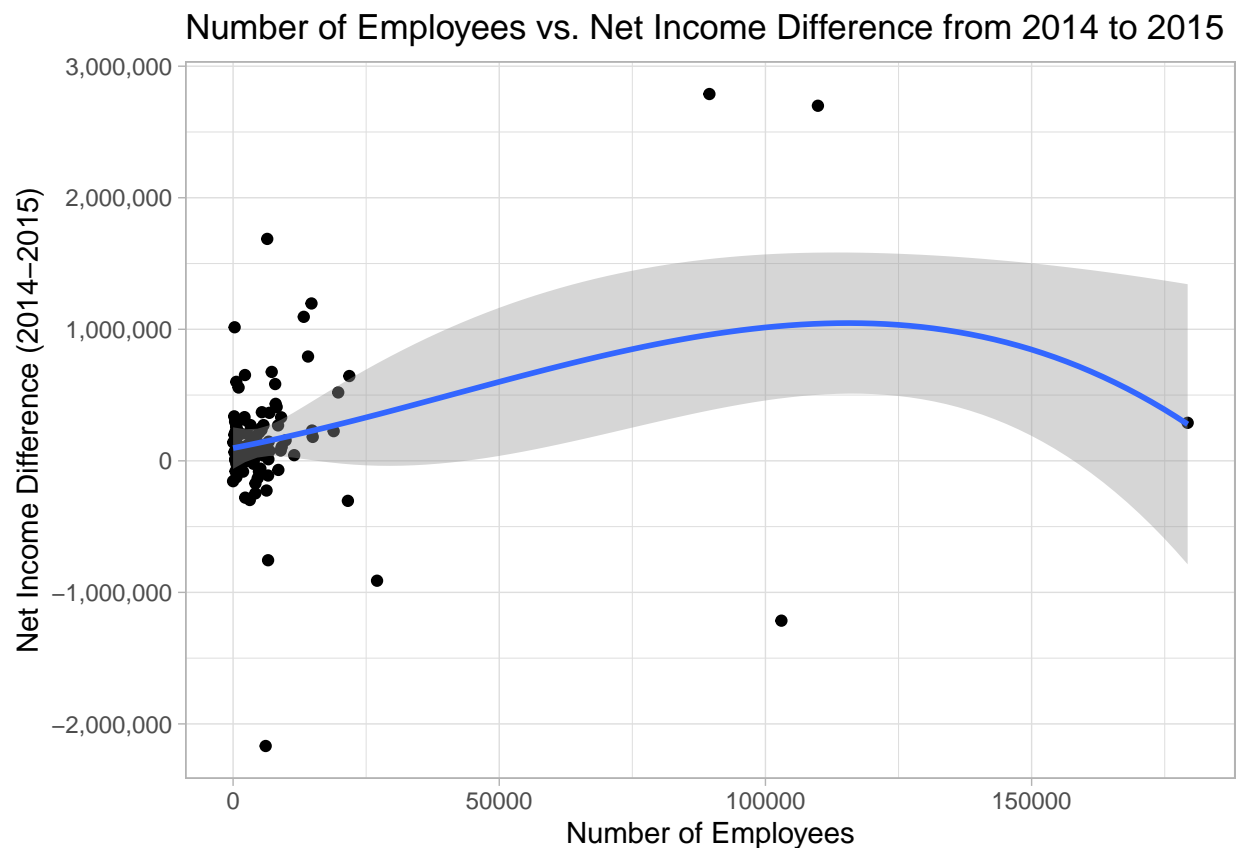
Create a new variable, `net_income_diff`, as the difference between the 2015 net income and the 2014 net income:

```
c_data <- c_data %>%
  mutate(net_income_diff = net_income_2015 - net_income_2014)
head(c_data$net_income_diff)
```

```
## [1] -1215100 2699298 2788500 164392 1687083 288980
```

****Produce a scatterplot of the net income difference (on the vertical axis) against the number of employees (on the horizontal axis). Add a smooth curve to your scatter plot.**


```
ggplot(c_data, aes(x = n_empl, y = net_income_diff)) +  
  theme_light() + geom_point() + geom_smooth(span = 1) +  
  scale_y_continuous(label = comma) +  
  labs(x = "Number of Employees", y = "Net Income Difference (2014-2015)",  
       title = "Number of Employees vs. Net Income Difference from 2014 to 2015")
```



Using only the companies with a positive number of employees and a positive net income difference, modify the above plot to use a logarithmic scale on both axis:

```
filter(c_data, n_empl > 0 & net_income_diff > 0) %>%  
ggplot(aes(x = n_empl, y = net_income_diff)) +  
  theme_light() + geom_point() + geom_smooth(span = 1) +  
  scale_x_log10(label = comma) + scale_y_log10(label = comma) +  
  labs(x = "Number of Employees", y = "Net Income Difference (2014-2015)",  
       title = "Number of Employees vs. Net Income Difference from 2014 to 2015")
```

Number of Employees vs. Net Income Difference from 2014 to 2015



Using all of the companies, create a factor, `sector_new`, based on the variable `sector`, such that the levels of the factor are labelled catering, (corresponding to 1), hotels (corresponding to 2), distribution (corresponding to 3), and communications (corresponding to 4):

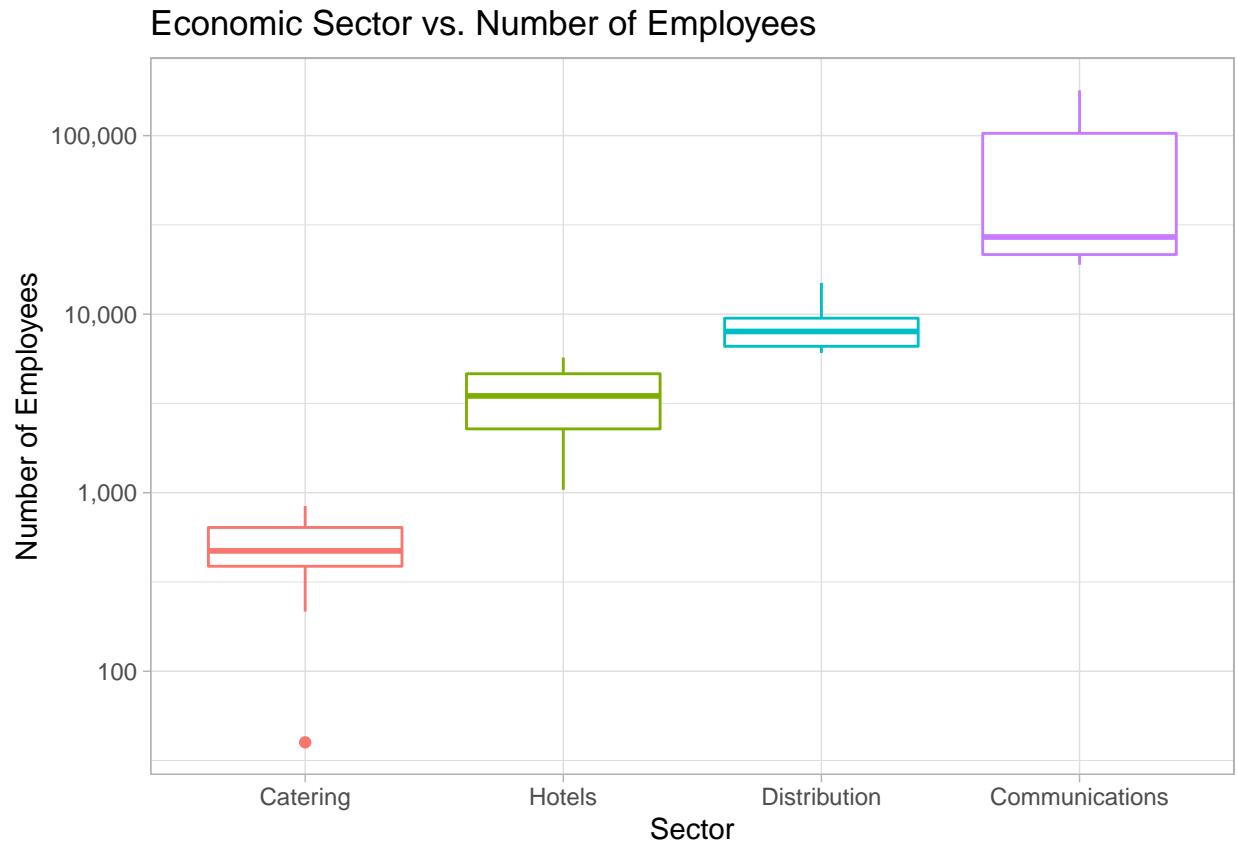
```
c_data <- c_data %>%
  mutate(sector_new = factor(sector, levels = 1:4,
                             labels = c('Catering', 'Hotels', 'Distribution',
                                         'Communications'))))

head(c_data$sector_new)
```

```
## [1] Communications Communications Communications Hotels          Distribution
## [6] Communications
## Levels: Catering Hotels Distribution Communications
```

Using all of the companies, produce a boxplot of `n_empl` against `sector_new`:

```
ggplot(c_data, aes(x = sector_new, y = n_empl, colour = sector_new)) +
  theme_light() + geom_boxplot() +
  scale_y_log10(label = comma) +
  labs(x = "Sector", y = "Number of Employees",
       title = "Economic Sector vs. Number of Employees") +
  theme(legend.position = 'none')
```



Using all of the companies, convert the continuous variable, `oper_result`, (which has units of GBP) into the continuous variable `oper_result_millions` (with units of millions of GBP):

```
c_data <- c_data %>%
  mutate(oper_result_millions = oper_result / 1000000)
head(c_data$oper_result_millions)
```

```
## [1] -0.277747  2.861590  2.747100  0.256078  2.025679 -4.356350
```


Using all of the companies, convert the continuous variable `oper_result_millions` into the factor `oper_result_new`, with breaks at -5, 0, 0.25, and 3, in such a way that the corresponding intervals are `[-5, 0)`, `[0, 0.25)` and `[0.25, 3)`:

```
c_data <- c_data %>%
  mutate(oper_result_new = cut(x = oper_result_millions, breaks = c(-5, 0, 0.25, 3)))
table(c_data$oper_result_new)
```

```
##
##  (-5,0] (0,0.25] (0.25,3]
##      23      65      12
```

Using all of the companies, produce histograms of `n_empl` split by `oper_result_new`, using a logarithmic scale for the number of employees:


```
ggplot(c_data, aes(x = n_empl, fill = oper_result_new)) +
  theme_light() + geom_histogram() +
  scale_x_log10(label = comma) +
  facet_grid(oper_result_new ~ .) +
  labs(x = "Number of Employees", y = "Count", fill = "Operational Result (millions £)",
       title = "Number of Employees by Operational Result")
```



Exercise 4: Displaying Student Assessment Data

The file `Module_Marks_Invented_Example.csv` contains some student assessment data. It would not be ethically correct to release students' marks, so these marks are entirely fabricated, although they do share some of the properties of real marks.

Students take three modules: M1, M2, and M3. The coursework/examination/overall marks are indicated by the suffix of the column title (*.C/.E/.F respectively*).

```
m_data <- read_csv('../data/Module_Marks_Invented_Example.csv')
head(m_data)
```

```
## # A tibble: 6 x 9
##   M1.C M1.E M1.F M2.C M2.E M2.F M3.C M3.E M3.F
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1    NA    NA    NA    NA    NA    NA    NA    NA    NA
## 2    NA    NA    NA    NA    NA    NA    NA    NA    NA
```

```
## 3    NA    NA    NA    66    74    72    72    49    56
## 4    NA    NA    NA    NA    NA    NA    NA    NA    NA
## 5    NA    NA    NA    NA    NA    NA    NA    NA    NA
## 6    NA    NA    NA    NA    NA    NA    NA    NA    NA
```

Note that as not all students take all modules, there are a lot of missing values.

Use the `gather()` function from the `tidyr` package to place all of the marks into one column names `marks`, with another column named `source` indicating the source of the marks (*i.e.* `M1.C`, `M1.E`, `M1.F`, `M2.C`, `M2.E`, `M2.F`, `M3.C`, `M3.E`, `M3.F`):

```
m_long <- gather(m_data, 'source', 'marks')
head(m_long)
```

```
## # A tibble: 6 x 2
##   source marks
##   <chr>   <dbl>
## 1 M1.C     NA
## 2 M1.C     NA
## 3 M1.C     NA
## 4 M1.C     NA
## 5 M1.C     NA
## 6 M1.C     NA
```

Use the `separate()` function from the `tidyr` package to separate the `source` column of `m_long` into two columns named `module` (containing the module code `M1`, `M2`, or `M3`) and `component` (containing the assessment component `C`, `E`, or `F`). Do not change the `marks` column:

```
m_long_2 <- m_long %>%
  separate(col = source, into = c('module', 'component'), sep = '\\.')
head(m_long_2)
```

```
## # A tibble: 6 x 3
##   module component marks
##   <chr>   <chr>     <dbl>
## 1 M1      C           NA
## 2 M1      C           NA
## 3 M1      C           NA
## 4 M1      C           NA
## 5 M1      C           NA
## 6 M1      C           NA
```

Use the `mutate()` function from `dplyr` to define `component` as a factor, named `component_f`, with levels: “C”, “E”, and “F” and labels: “Coursework”, “Examination”, and “Overall”:

```
m_long_3 <- m_long_2 %>%
  mutate(component_f = factor(component, levels = c('C', 'E', 'F'),
                                labels = c('Coursework', 'Examination', 'Overall')))
head(m_long_3$component_f)
```

```
## [1] Coursework Coursework Coursework Coursework Coursework Coursework
## Levels: Coursework Examination Overall
```

Finally, produce a boxplot of this data using ggplot2:

```
ggplot(m_long_3, aes(x = component_f, y = marks, fill = component_f)) +
  theme_light() + geom_boxplot() + facet_grid(. ~ module) +
  scale_fill_manual(values = c('orange', 'yellow', 'magenta')) +
  scale_y_continuous(breaks = c(30, 40, 50, 60, 70, 100),
                    minor_breaks = c(30, 40, 50, 60, 70, 80, 90, 100),
                    limits = c(30, 100)) +
  labs(x = "Component", y = "Mark (%)") +
  theme(axis.text.x = element_text(size = 14, angle = 90, vjust = 0.5, hjust = 1),
        legend.position = 'none')
```

