

Session 13 - Summary Statistics, Linear Model, and Correlation

Alex Mounsey

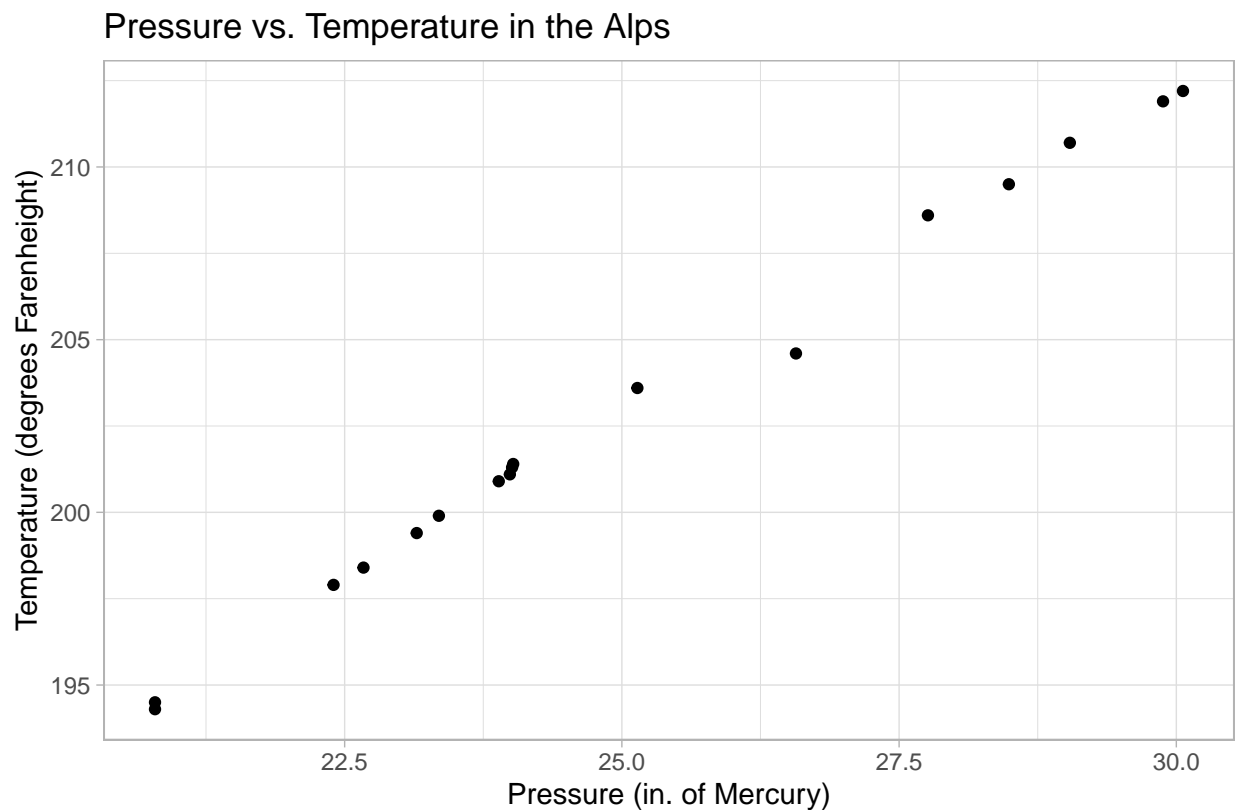
16/11/2020

Exercise 1: Simple Linear Regression

Consider the famous and scientifically important `Forbes` data. This data comprises of 17 pairs of numbers corresponding to the observed boiling point and corrected barometric pressure at locations in the Alps.

Use the `ggplot2` package to plot the data in the `Forbes` dataframe, placing the `pres` variable on the horizontal axis and the `bp` variable on the vertical axis:

```
ggplot(Forbes, aes(x = pres, y = bp)) +  
  theme_light() + geom_point() +  
  labs(x = "Pressure (in. of Mercury)", y = "Temperature (degrees Farenheight)",  
        title = "Pressure vs. Temperature in the Alps",  
        caption = "Source: Forbes data via ALR4 R Package")
```



Source: Forbes data via ALR4 R Package

Calculate the mean, median, variance, standard deviation, and interquartile range (*IQR*) for both pres and bp. Do this using dplyr and again using the with() function:

```
# Solution with dplyr
Forbes %>%
  summarise(mean_pres = mean(pres), median_pres = median(pres), var_pres = var(pres),
            sd_pres = sd(pres), iqr_pres = IQR(pres), mean_bp = mean(bp),
            median_bp = median(bp), var_bp = var(bp), sd_bp = sd(bp), iqr_bp = IQR(bp))
```

```
##   mean_pres median_pres var_pres sd_pres iqr_pres mean_bp median_bp var_bp
## 1  25.05882    24.01 9.121111 3.020118    4.61 202.9529    201.3 33.1739
##      sd_bp iqr_bp
## 1 5.759679    9.2
```

```
# Solution with 'with()'
mean_pres <- with(Forbes, mean(pres))
```

Fit the Simple Linear Regression Model:

```
m <- lm(bp ~ pres, data = Forbes)
coef(m)
```

```
## (Intercept)      pres
## 155.296483    1.901784
```


Write down a mathematical expression for the model:

The estimate for α is $\hat{\alpha} = 155.3$ and the estimate for β is $\hat{\beta} = 1.902$. The fitted model is therefore $Temperature = \hat{\alpha} + \hat{\beta}Pressure = 155.3 + 1.902Pressure$.

Use the geom_smooth() function from ggplot2 to show, in blue, the fitted line on the plot created above:

```
ggplot(Forbes, aes(x = pres, y = bp)) +
  theme_light() + geom_point() +
  geom_smooth(method = 'lm', se = F, colour = 'blue') +
  labs(x = "Pressure (in. of Mercury)", y = "Temperature (degrees Farenheight)",
       title = "Pressure vs. Temperature in the Alps",
       caption = "Source: Forbes data via ALR4 R package")
```

Pressure vs. Temperature in the Alps



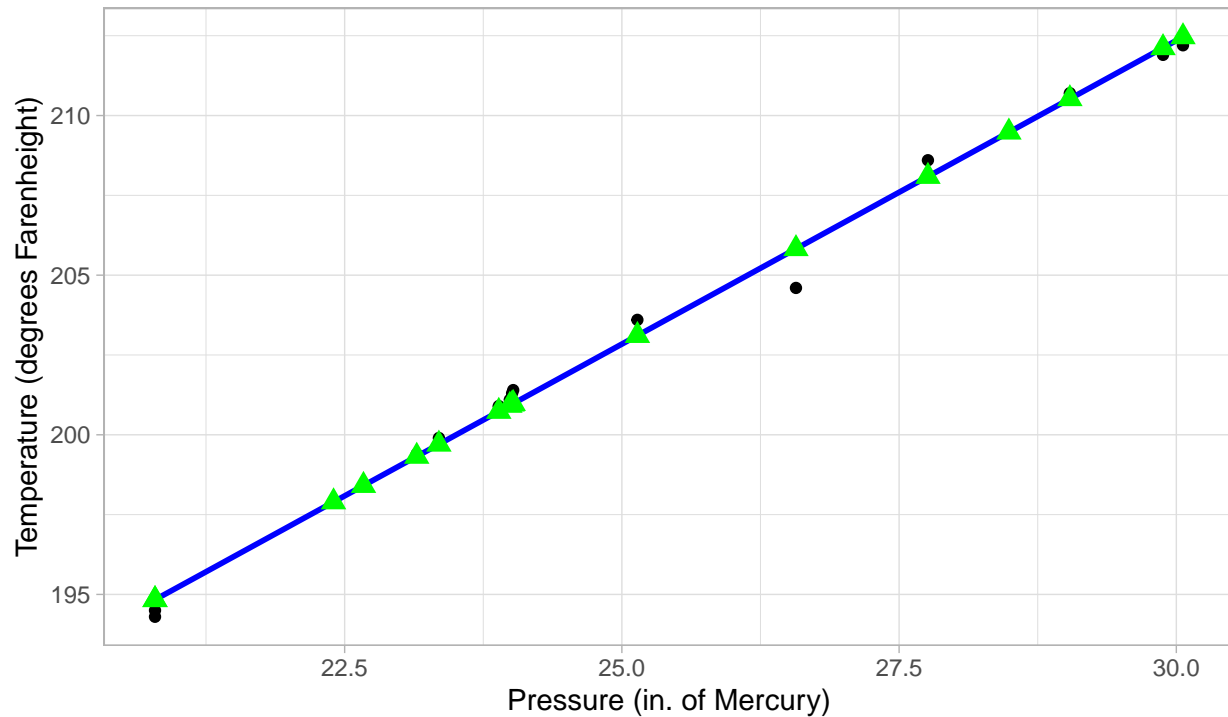
Source: Forbes data via ALR4 R package

Use the function `add_predictions()` from `modelr` to compute the fitted values and include them with the original data. Add these values to your plot in green:

```
Forbes %>%
  add_predictions(m) %>%
  ggplot(aes(x = pres, y = bp)) +
  theme_light() + geom_point() +
  geom_smooth(method = 'lm', se = F, colour = 'blue') +
  geom_point(aes(y = pred), col = 'green', pch = 17, size = 3) +
  labs(x = "Pressure (in. of Mercury)", y = "Temperature (degrees Fahrenheit)",
       title = "Pressure vs. Temperature in the Alps",
       subtitle = "Fitted values shown in green",
       caption = "Source: Forbes data via ALR4 R package")
```

Pressure vs. Temperature in the Alps

Fitted values shown in green



Source: Forbes data via ALR4 R package

Is there an underlying relationship between temperature and pressure? Justify your conclusion:

```
p_value <- signif(summary(m)$coefficients[2,4], 3)
summary(m)
```

```
##
## Call:
## lm(formula = bp ~ pres, data = Forbes)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.22687 -0.22178  0.07723  0.19687  0.51001
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 155.29648    0.92734   167.47  <2e-16 ***
## pres         1.90178     0.03676    51.74  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.444 on 15 degrees of freedom
## Multiple R-squared:  0.9944, Adjusted R-squared:  0.9941
## F-statistic: 2677 on 1 and 15 DF, p-value: < 2.2e-16
```

The p-value is equal to 2.53×10^{-18} , which is less than 0.05 - therefore we reject the null hypothesis

$H_0 : \beta = 0$ (there is no underlying relationship) in favour of the alternative hypothesis $H_1 : \beta \neq 0$ (there is an underlying relationship).

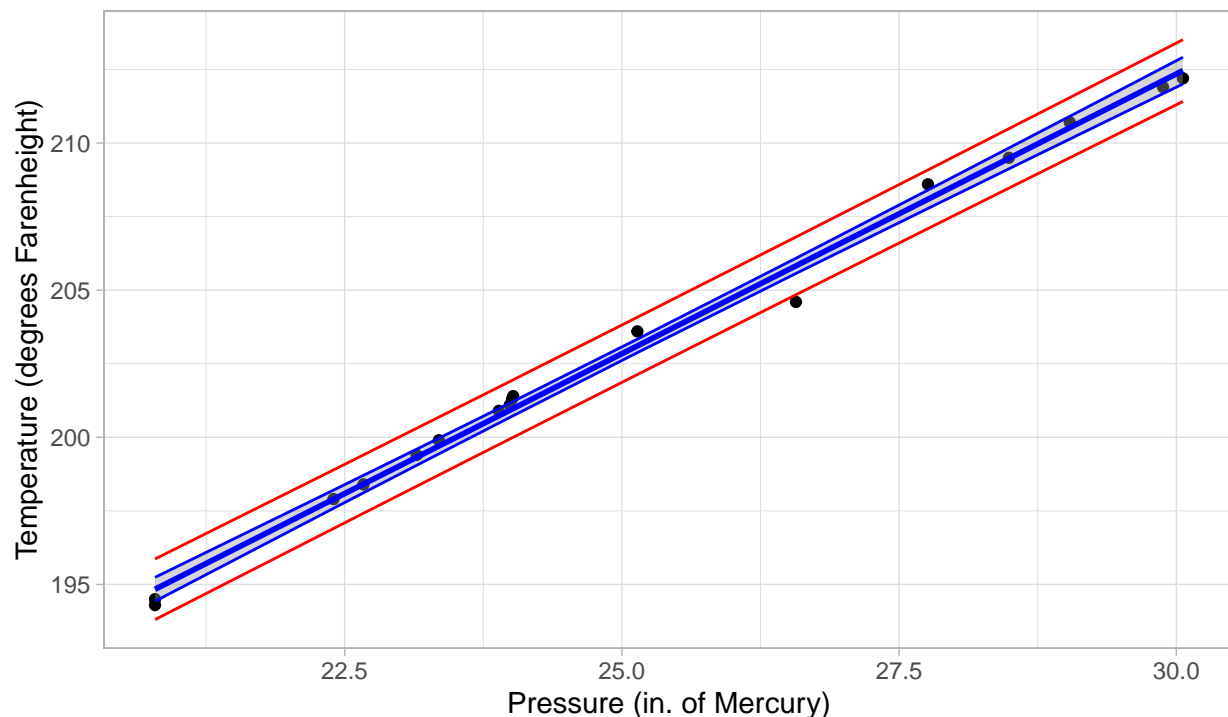
Add the confidence and prediction intervals to your graph:

```
confidence <- data.frame(Forbes, predict(m, interval = 'confidence'))
prediction <- data.frame(Forbes, predict(m, interval = 'prediction'))

Forbes %>%
  add_predictions(m) %>%
  ggplot(aes(x = pres, y = bp)) +
  theme_light() + geom_point() +
  geom_smooth(method = 'lm', se = T, colour = 'blue') +
  geom_line(data = confidence, aes(y = lwr), colour = 'blue') +
  geom_line(data = confidence, aes(y = upr), colour = 'blue') +
  geom_line(data = prediction, aes(y = lwr), colour = 'red') +
  geom_line(data = prediction, aes(y = upr), colour = 'red') +
  labs(x = "Pressure (in. of Mercury)", y = "Temperature (degrees Farenheight)",
       title = "Pressure vs. Temperature in the Alps",
       subtitle = "with Confidence intervals (blue) and Prediction intervals (red)",
       caption = "Source: Forbes data via ALR4 R package")
```

Pressure vs. Temperature in the Alps

with Confidence intervals (blue) and Prediction intervals (red)



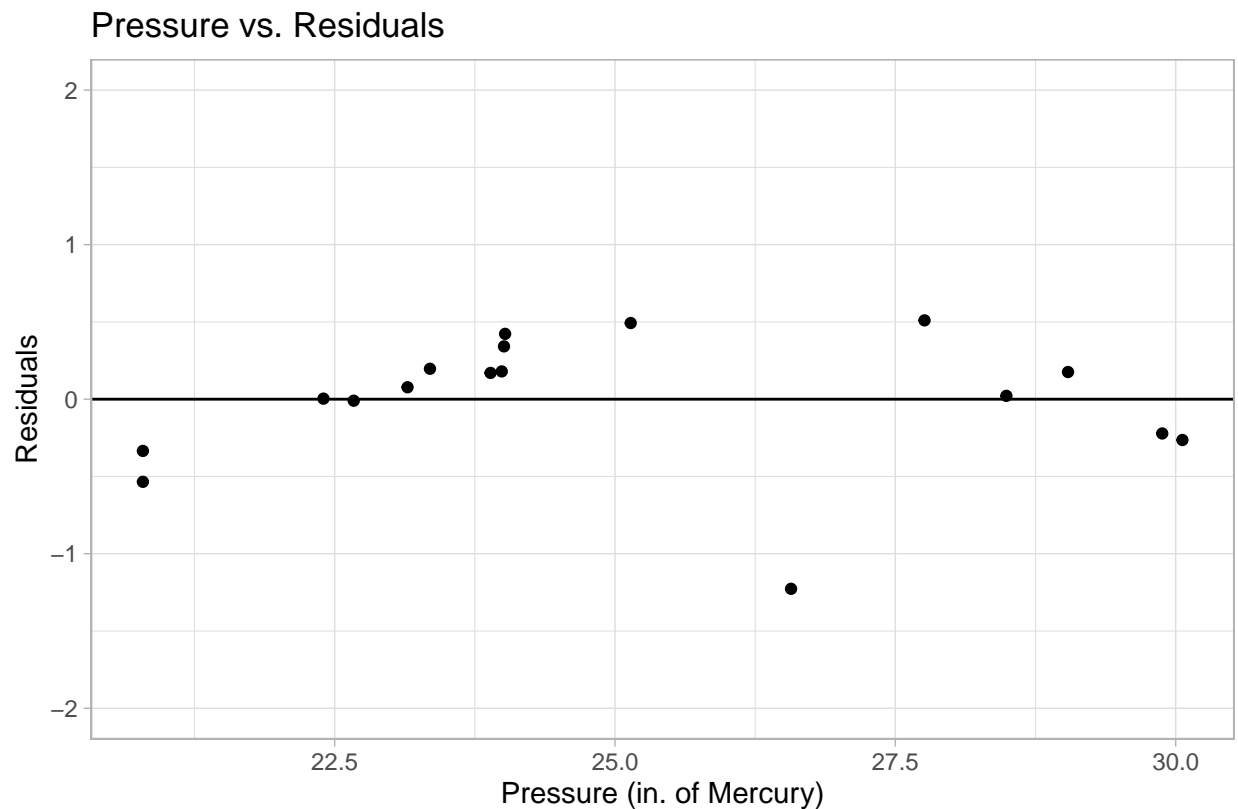
Source: Forbes data via ALR4 R package

Using the `add_residuals()` function from `modelr`, compute the residuals and include them with the original data and fitted values:

```

Forbes %>%
  add_predictions(m) %>%
  add_residuals(m) %>%
  ggplot(aes(x = pres, y = resid)) +
    theme_light() + geom_point() +
    scale_y_continuous(limits = c(-2, 2)) +
    geom_hline(aes(yintercept = 0)) +
    labs(x = "Pressure (in. of Mercury)", y = "Residuals",
         title = "Pressure vs. Residuals",
         caption = "Source: Forbes data via ALR4 R package")

```




```

Forbes %>%
  add_predictions(m) %>%
  add_residuals(m) %>%
  ggplot(aes(x = pred, y = resid)) +
    theme_light() + geom_point() +
    scale_y_continuous(limits = c(-2, 2)) +
    geom_hline(aes(yintercept = 0)) +
    labs(x = "Fitted Values", y = "Residuals",
         title = "Fitted Values vs. Residuals",
         caption = "Source: Forbes data via ALR4 R package")

```

Fitted Values vs. Residuals



Source: Forbes data via ALR4 R package

What is the correlation coefficient between temperature and pressure? Use a test of correlation to determine whether there is an underlying dependency between temperature and pressure:

```
p_value <- with(Forbes, cor(pres, bp))
with(Forbes, cor.test(pres, bp))

##
## Pearson's product-moment correlation
##
## data:  pres and bp
## t = 51.741, df = 15, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
##  0.9920669 0.9990205
## sample estimates:
##      cor
## 0.9972102
```

Once again, the p-value is below 0.05 (0.9972102). This prompts us to reject the null hypothesis $H_0 : \rho = 0$ in favor of the alternative hypothesis $H_1 : \rho \neq 0$. Therefore, we can conclude that there *is* an underlying linear dependency between temperature and pressure.

Exercise 2: Multiple Regression

The aim of this exercise is to extend the simple linear regression model:

$$y = \alpha + \beta x + error$$

which has one explanatory variable (x) to the multiple linear regression model:

$$y = \alpha + \beta_1 x_1 + \beta_2 x_2 + error$$

which has more than one explanatory variable (here: x_1 and x_2).

Here is a dataset which will help us understand the multiple linear regression model. The data comes from shops in ten towns:

Town	Advertising (£000s)	Population (000s)	Annual Sales (£000s)
1	2	41	58
2	6	93	105
3	8	78	88
4	8	102	118
5	12	96	117
6	16	112	137
7	20	119	157
8	20	131	169
9	22	118	149
10	26	185	202

Input this data:


```
town <- 1:10
advertising <- c(2, 6, 8, 8, 12, 16, 20, 20, 22, 26)
population <- c(41, 93, 78, 102, 96, 112, 119, 131, 118, 185)
sales <- c(58, 105, 88, 118, 117, 137, 157, 169, 149, 202)

shops <- data.frame(town, advertising, population, sales)
```

Plot the annual sales against advertising, illustrating the dependence of sales on advertising, using a smooth curve:

```
ggplot(shops, aes(x = advertising, y = sales)) +
  theme_light() + geom_point() + geom_smooth() +
  labs(x = "Advertising (£'000s)", y = "Sales Revenue (£'000s)",
       title = "Advertising Budget vs. Sales Revenue")
```



Advertising Budget vs. Sales Revenue



Plot annual sales against population, illustrating the dependence of sales on population, using a smooth curve:

```
ggplot(shops, aes(x = population, y = sales)) +  
  theme_light() + geom_point() + geom_smooth() +  
  labs(x = "Population (000s)", y = "Sales Revenue (£'000s)",  
       title = "Population vs. Sales Revenue")
```

Population vs. Sales Revenue



In light of the linear relationships observed above, we can fit the model:

$$sales = \alpha + \beta_1 advertising + \beta_2 population + error$$

Which can be implemented in R with the following code:

```
m <- lm(sales ~ advertising + population, data = shops)
p_value <- signif(summary(m)$coefficients[2,4], 3)
summary(m)
```

```
##
## Call:
## lm(formula = sales ~ advertising + population, data = shops)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.9777 -3.9884 -0.5326  3.7399 10.1376
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   26.1393     8.0014   3.267  0.01373 *
## advertising    2.0949     0.6338   3.305  0.01302 *
## population     0.6933     0.1352   5.129  0.00135 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 6.777 on 7 degrees of freedom
## Multiple R-squared:  0.9796, Adjusted R-squared:  0.9737
## F-statistic: 167.7 on 2 and 7 DF,  p-value: 1.221e-06
```

The second p-value (0.013) can be used to check whether β_1 is significantly different from zero. If this p-value is less than 0.05, we can conclude that advertising does have an effect on sales for all similar shops.

We can make predictions in the usual way; for example, to predict sales if £10,000 is spent on advertising in a town with a population of 140,000:

```
predict(m, newdata = data.frame(advertising = 10, population = 140))
```

```
##          1
## 144.1537
```

Produce a prediction for sales, together with a confidence and prediction interval, if £10,000 is spent on advertising in a town of 140,000 people and if £30,000 is spent on advertising in a town of 200,000 people:

```
# Confidence intervals
predict(m, newdata = data.frame(advertising = c(10, 30), population = c(140, 200)),
       interval = 'confidence')
```

```
##          fit      lwr      upr
## 1 144.1537 127.3959 160.9115
## 2 227.6505 213.2405 242.0606
```

```
# Prediction intervals
predict(m, newdata = data.frame(advertising = c(10, 30), population = c(140, 200)),
       interval = 'prediction')
```

```
##          fit      lwr      upr
## 1 144.1537 120.9663 167.3411
## 2 227.6505 206.0987 249.2024
```

Exercise 3: Multiple Regression

The data below shows the loading times of cargo in relation to weight, volume, and the number of items in the shipment:

Cargo	Weight	Volume	No. of Items	Loading Time
1	0.4	53	158	64
2	0.4	23	163	60
3	3.1	19	37	71
4	0.6	34	157	61
5	4.7	24	59	54
6	1.7	65	123	77
7	9.4	44	46	81
8	10.1	31	117	93
9	11.6	29	173	93

Cargo	Weight	Volume	No. of Items	Loading Time
10	12.6	58	112	51
11	10.9	37	111	76
12	23.1	46	114	96
13	23.1	50	134	77
14	21.6	44	73	93
15	23.1	56	168	95
16	1.9	36	143	54
17	26.8	58	202	168
18	29.9	51	124	99

Input the data:

```
cargo <- 1:18
weight <- c(0.4, 0.4, 3.1, 0.6, 4.7, 1.7, 9.4, 10.1, 11.6, 12.6, 10.9, 23.1, 23.1, 21.6, 23.1, 1.9, 26.8)
volume <- c(53, 23, 19, 34, 24, 65, 44, 31, 29, 58, 37, 46, 50, 44, 56, 36, 58, 51)
n_items <- c(158, 163, 37, 157, 59, 123, 46, 117, 173, 112, 111, 114, 134, 73, 168, 143, 202, 124)
loading_time <- c(64, 60, 71, 61, 54, 77, 81, 93, 93, 51, 76, 96, 77, 93, 95, 54, 168, 99)

shipping <- data.frame(cargo, weight, volume, n_items, loading_time)
```

Fit the model:

$$\text{LoadingTime} = \alpha\beta_1\text{Weight} + \beta_2\text{Volume} + \beta_3\text{NumberOfItems} + \text{error}$$

```
m <- lm(loading_time ~ weight + volume + n_items, data = shipping)
coef(m)
```

```
## (Intercept)      weight      volume      n_items
## 43.65219779  1.78477968 -0.08339706  0.16113269
```

Which variable has the highest p-value?

```
summary(m)
```

```
##
## Call:
## lm(formula = loading_time ~ weight + volume + n_items, data = shipping)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -28.35 -11.39  -2.66   12.09   48.80
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  43.6522    18.0102   2.424  0.02949 *
## weight       1.7848     0.5377   3.319  0.00506 **
## volume      -0.0834     0.4177  -0.200  0.84462
## n_items       0.1611     0.1117   1.443  0.17102
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.97 on 14 degrees of freedom
## Multiple R-squared:  0.5493, Adjusted R-squared:  0.4528
## F-statistic: 5.689 on 3 and 14 DF,  p-value: 0.009224
```

Remove the variable with the highest p-value from the model, provided that the p-value is greater than 0.05:

```
m <- lm(loading_time ~ weight + n_items, data = shipping)
coef(m)
```

```
## (Intercept)      weight      n_items
##   41.479364    1.737438    0.154843
```

Which variable has the highest p-value?

```
summary(m)
```

```
##
## Call:
## lm(formula = loading_time ~ weight + n_items, data = shipping)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -29.713 -11.324  -2.953   11.286   48.679
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   41.4794    13.8834   2.988  0.00920 **
## weight         1.7374     0.4669   3.721  0.00205 **
## n_items        0.1548     0.1036   1.494  0.15592
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 19.32 on 15 degrees of freedom
## Multiple R-squared:  0.5481, Adjusted R-squared:  0.4878
## F-statistic: 9.095 on 2 and 15 DF,  p-value: 0.002589
```

Remove the variable with the highest p-value from the model, provided that the p-value is greater than 0.05:

```
m <- lm(loading_time ~ weight, data = shipping)
coef(m)
```

```
## (Intercept)      weight
##   59.258959    1.843436
```

Does weight have an effect on loading time for all cargo?

```
summary(m)
```

```
##
## Call:
## lm(formula = loading_time ~ weight, data = shipping)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -31.486  -8.282  -1.674   5.623  59.337
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   59.2590     7.4200   7.986 5.67e-07 ***
## weight        1.8434     0.4789   3.849 0.00142 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 20.05 on 16 degrees of freedom
## Multiple R-squared:  0.4808, Adjusted R-squared:  0.4484
## F-statistic: 14.82 on 1 and 16 DF,  p-value: 0.001417
```

The procedure above is called *backward elimination*, and has simplified the original model:

$$LoadingTime = \alpha + \beta_1 Weight + \beta_2 Volume + \beta_3 NumberOfItems + error$$


to:

$$LoadingTime = \alpha + \beta Weight + error$$

Use `ggplot2` to plot `weight` vs. `loading_time`, with the cargo number as the plotting characters:

```
ggplot(shipping, aes(x = weight, y = loading_time, label = cargo)) +
  theme_light() + geom_text() +
  geom_smooth(method = 'lm', se = FALSE) +
  labs(x = "Weight", y = "Loading Time",
       title = "Cargo Weight vs. Loading Time")
```


Cargo Weight vs. Loading Time



Produce residual plots:


```
shipping %>%
  add_predictions(m) %>%
  add_residuals(m) %>%
  ggplot(aes(x = weight, y = resid, label = cargo)) +
    theme_light() + geom_text() +
    scale_y_continuous(limits = c(-70, 70)) +
    geom_hline(aes(yintercept = 0)) +
    labs(x = "Weight", y = "Residuals",
         title = "Weight vs. Residuals")
```

Weight vs. Residuals



```
shipping %>%
  add_predictions(m) %>%
  add_residuals(m) %>%
  ggplot(aes(x = pred, y = resid, label = cargo)) +
    theme_light() + geom_text() +
    scale_y_continuous(limits = c(-70, 70)) +
    geom_hline(aes(yintercept = 0)) +
    labs(x = "Fitted Values", y = "Residuals",
         title = "Fitted Values vs. Residuals")
```


Fitted Values vs. Residuals



Comment on how well the simple linear regression model fits the data:


The residuals appear to be evenly scattered about 0, with no distinct pattern. This indicates that the variance is constant and the simple linear regression fits the data well.

Add, in red, the simple linear regression line based on all data except the 17th data point (*the outlier*):

```
ggplot(shipping, aes(x = weight, y = loading_time)) +
  theme_light() + geom_point() +
  geom_smooth(method = 'lm', se = F) +
  geom_smooth(data = shipping[-17,], method = 'lm', se = F, colour = 'red') +
  labs(x = "Weight", y = "Loading Time",
       title = "Cargo Weight vs. Loading Time",
       subtitle = "Linear model marked in red ignores outliers")
```

Cargo Weight vs. Loading Time

Linear model marked in red ignores outliers



What is the effect of removing the 17th data point?

The linear model which excludes the 17th data point in its calculation appears to fit the data more appropriately, and should produce more accurate predictions.


Exercise 4: Transformations

In previous sessions, we have worked with the questionnaire data. In particular, we used this code to produce the following plot:

```
q_data <- read_csv('../data/MATH513_Questionnaire_Data.csv')

ggplot(q_data, aes(x = Distance, y = Travel_time)) +
  theme_light() + geom_point() + geom_smooth() +
  geom_smooth(method = 'lm', se = F, colour = 'red') +
  scale_x_log10() + scale_y_log10() +
  labs(x = "Distance (miles)", y = "Time (minutes)",
       title = "Distance vs. Time")
```

Distance vs. Time

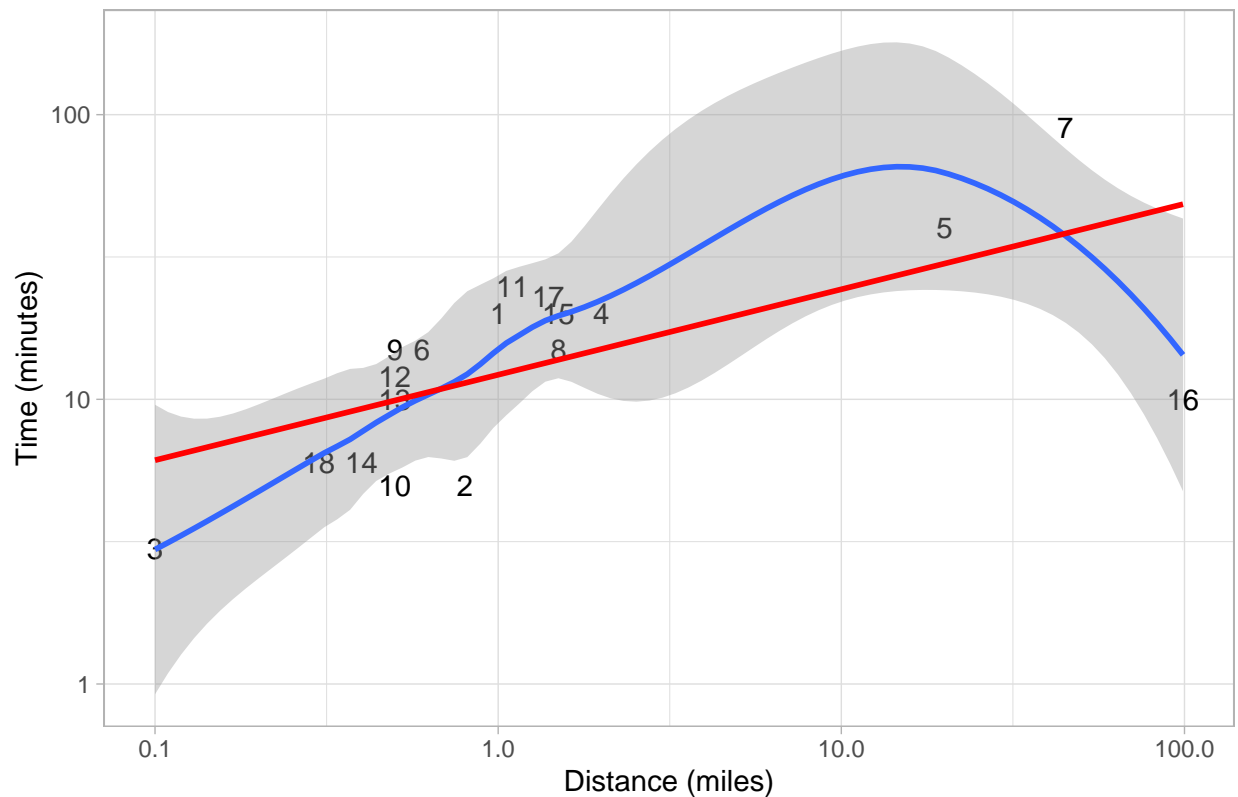


Add labels for each of the data points' index, and use them as plotting characters:

```
q_data <- q_data %>%
  mutate(data_point = 1:nrow(q_data))

ggplot(q_data, aes(x = Distance, y = Travel_time, label = data_point)) +
  theme_light() + geom_text() + geom_smooth() +
  geom_smooth(method = 'lm', se = F, colour = 'red') +
  scale_x_log10() + scale_y_log10() +
  labs(x = "Distance (miles)", y = "Time (minutes)",
       title = "Distance vs. Time")
```

Distance vs. Time



Fit the model:

$$\log(\text{TravelTime}) = \beta_0 + \beta_1 \log(\text{Distance}) + \text{error}$$

Note that as it's not possible to compute the log of zero, we add a small value (0.01) to Distance in the case that there are some zero distances.

```
m_log <- lm(log(Travel_time) ~ log(Distance + 0.01), data = q_data)
summary(m_log)
```

```
##
## Call:
## lm(formula = log(Travel_time) ~ log(Distance + 0.01), data = q_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.5771 -0.4146  0.2407  0.4051  0.8575
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.49588    0.15726  15.871 3.27e-11 ***
## log(Distance + 0.01)  0.30115    0.09023   3.338 0.00417 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## Residual standard error: 0.6561 on 16 degrees of freedom
## Multiple R-squared:  0.4105, Adjusted R-squared:  0.3736
## F-statistic: 11.14 on 1 and 16 DF,  p-value: 0.004174
```

Re-fit the model, omitting the 16th data point:

```
m_log <- lm(log(Travel_time) ~ log(Distance + 0.01), data = q_data, subset = -16)
```

Use this model to produce confidence intervals for β_0 and β_1 using the `cofint()` function:

```
cofint(m_log)
```

```
##                2.5 %   97.5 %
## (Intercept)      2.353922 2.795413
## log(Distance + 0.01) 0.357343 0.671532
```

These intervals give us an indication of the reliability of the estimates for β_0 and β_1 . The wider the interval, the less reliable the estimate is.

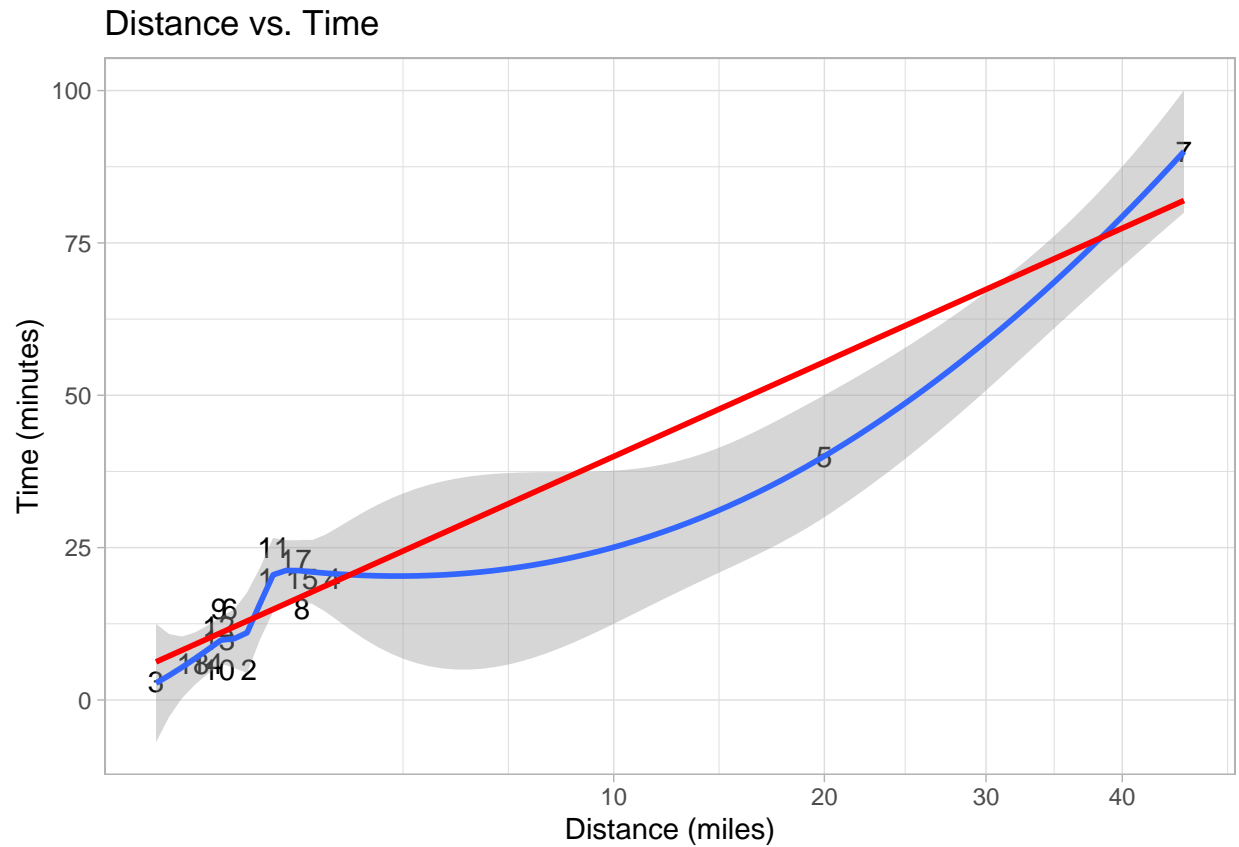
A mathematical argument that we do not discuss here tells us that the second of these confidence intervals, here $(0.357343, 0.671532)$, contains 0.5, then the travel time may depend on the square root of the distance.

To check this out, we should plot the data using a square root scale on the x-axis, with a standard linear scale on the y-axis. A square root scale spreads out values less than 1 and squashes values greater than 1.

In `ggplot2`, a square root scale on the x-axis can be produced using `scale_x_sqrt()` instead of `scale_x_log10()`.

Produce a plot which uses a square root scale on the x-axis and a standard linear scale on the y-axis:

```
ggplot(q_data[-16,], aes(x = Distance, y = Travel_time, label = data_point)) +
  theme_light() + geom_text() + geom_smooth() +
  geom_smooth(method = 'lm', se = F, col = 'red') +
  scale_x_sqrt() +
  labs(x = "Distance (miles)", y = "Time (minutes)",
       title = "Distance vs. Time")
```

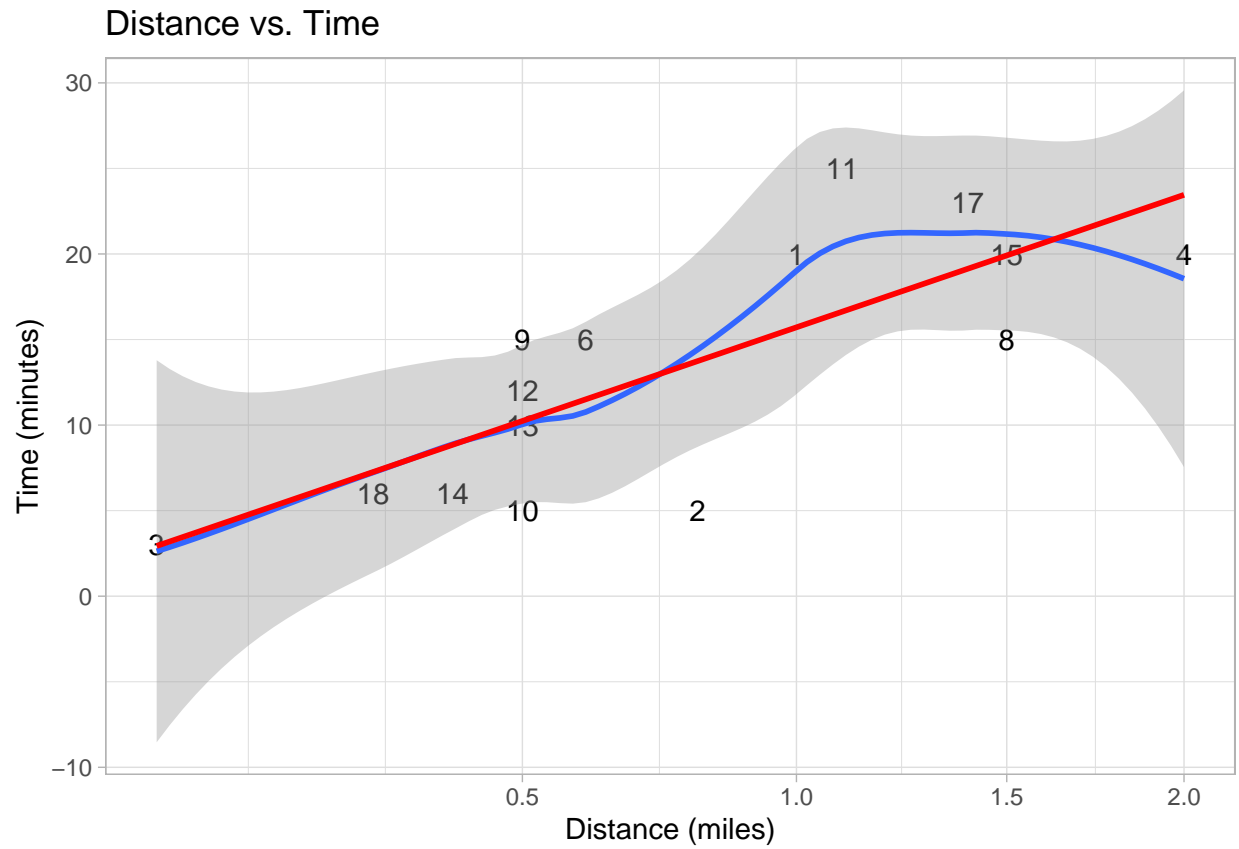


Comment on this plot:

All of the data points, excluding the 5th and 7th, are bunched together.

Produce the same plot, omitting the 5th and 7th data points:

```
ggplot(q_data[-c(5, 7, 16),], aes(x = Distance, y = Travel_time, label = data_point)) +
  theme_light() + geom_text() + geom_smooth() +
  geom_smooth(method = 'lm', se = F, col = 'red') +
  scale_x_sqrt() +
  labs(x = "Distance (miles)", y = "Time (minutes)",
       title = "Distance vs. Time")
```



Comment on this plot:

The travel time appears to depend on the square root of the distance.